

XGBoost



```
print('XGBoost model - SMOTE RFM')
```

```
xgb_model = xgb.XGBClassifier(objective='binary:logistic', eval_metric='auc',  
    learning_rate =0.3,    #change 0.01 >> 0.3  
    n_estimators=100,  
    max_depth=20,    #change 2 >> 20  
    gamma=0.0,  
    colsample_bytree=0.6)
```

```
predicted_y = []  
expected_y = []
```

```
xgb_model_SMOTE_rfm = xgb_model.fit(X_SMOTE_rfm, y_SMOTE_rfm, early_stopping_rounds=5, eval_set=[(X_test_rfm.to_numpy(), y_test_rfm)])  
predictions = xgb_model_SMOTE_rfm.predict(X_SMOTE_rfm)  
predicted_y.extend(predictions)  
expected_y.extend(y_SMOTE_rfm)  
report_train = classification_report(expected_y, predicted_y)  
print('training set')  
print(report_train)
```

```
predicted_y = []  
expected_y = []  
predictions = xgb_model_SMOTE_rfm.predict(X_test_rfm.to_numpy())  
predicted_y.extend(predictions)  
expected_y.extend(y_test_rfm)  
report_test = classification_report(expected_y, predicted_y)  
print('test set')  
print(report_test)
```



XGBoost model - SMOTE RFM

[0] validation_0-auc:0.562636

Will train until validation_0-auc hasn't improved in 5 rounds.

[1] validation_0-auc:0.636721

[2] validation_0-auc:0.639358

[3] validation_0-auc:0.644874

[4] validation_0-auc:0.663385

[5] validation_0-auc:0.667483

[6] validation_0-auc:0.662065

[7] validation_0-auc:0.674141

[8] validation_0-auc:0.681083

[9] validation_0-auc:0.681651

[10] validation_0-auc:0.68425

[11] validation_0-auc:0.685515

[12] validation_0-auc:0.687809

[13] validation_0-auc:0.692853

[14] validation_0-auc:0.697338

[15] validation_0-auc:0.701752

[16] validation_0-auc:0.705411

[17] validation_0-auc:0.708047

[18] validation_0-auc:0.710109

[19] validation_0-auc:0.711633

[20] validation_0-auc:0.711728

[21] validation_0-auc:0.712254

[22] validation_0-auc:0.708893

[23] validation_0-auc:0.708511

[24] validation_0-auc:0.709084

[25] validation_0-auc:0.709794

[26] validation_0-auc:0.710745

Stopping. Best iteration:

[21] validation_0-auc:0.712254

training set

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.92	0.99	0.95	4389
---	------	------	------	------

1	0.98	0.92	0.95	4389
---	------	------	------	------

accuracy			0.95	8778
----------	--	--	------	------

macro avg	0.95	0.95	0.95	8778
-----------	------	------	------	------

weighted avg	0.95	0.95	0.95	8778
--------------	------	------	------	------

test set

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.90	0.96	0.93	1848
---	------	------	------	------

1	0.23	0.10	0.14	218
---	------	------	------	-----

accuracy			0.87	2066
----------	--	--	------	------

macro avg	0.56	0.53	0.53	2066
-----------	------	------	------	------

weighted avg	0.83	0.87	0.85	2066
--------------	------	------	------	------



```
print('XGBoost model - SMOTE CLV')
```

```
xgb_model = xgb.XGBClassifier(objective='binary:logistic', eval_metric='auc',  
learning_rate =0.3,      #change 0.01 >> 0.3  
n_estimators=100,  
max_depth=20,           #change 2 >> 20  
gamma=0.0,  
colsample_bytree=0.6)
```

```
predicted_y = []  
expected_y = []
```

```
xgb_model_SMOTE_clv = xgb_model.fit(X_SMOTE_clv, y_SMOTE_clv, early_stopping_rounds=5, eval_set=[(X_test_clv.to_numpy(), y_test_clv)])  
predictions = xgb_model_SMOTE_clv.predict(X_SMOTE_clv)  
predicted_y.extend(predictions)  
expected_y.extend(y_SMOTE_clv)  
report_train = classification_report(expected_y, predicted_y)  
print('training set')  
print(report_train)
```

```
predicted_y = []  
expected_y = []  
predictions = xgb_model_SMOTE_clv.predict(X_test_clv.to_numpy())  
predicted_y.extend(predictions)  
expected_y.extend(y_test_clv)  
report_test = classification_report(expected_y, predicted_y)  
print('test set')  
print(report_test)
```



XGBoost model - SMOTE CLV

[0] validation_0-auc:0.616022

Will train until validation_0-auc hasn't improved in 5 rounds.

[1] validation_0-auc:0.623804

[2] validation_0-auc:0.627667

[3] validation_0-auc:0.63077

[4] validation_0-auc:0.634964

[5] validation_0-auc:0.633452

[6] validation_0-auc:0.643062

[7] validation_0-auc:0.648257

[8] validation_0-auc:0.652092

[9] validation_0-auc:0.654402

[10] validation_0-auc:0.655087

[11] validation_0-auc:0.654912

[12] validation_0-auc:0.652753

[13] validation_0-auc:0.650833

[14] validation_0-auc:0.651853

[15] validation_0-auc:0.654077

Stopping. Best iteration:

[10] validation_0-auc:0.655087

training set

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.98	0.99	0.99	4389
---	------	------	------	------

1	0.99	0.98	0.99	4389
---	------	------	------	------

accuracy			0.99	8778
----------	--	--	------	------

macro avg	0.99	0.99	0.99	8778
-----------	------	------	------	------

weighted avg	0.99	0.99	0.99	8778
--------------	------	------	------	------

test set

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.90	0.92	0.91	1848
---	------	------	------	------

1	0.19	0.16	0.17	218
---	------	------	------	-----

accuracy			0.84	2066
----------	--	--	------	------

macro avg	0.55	0.54	0.54	2066
-----------	------	------	------	------

weighted avg	0.83	0.84	0.83	2066
--------------	------	------	------	------



hyper parameter tuning - grid search

```
from sklearn.model_selection import KFold, GridSearchCV
from sklearn.metrics import accuracy_score, make_scorer
# Define our search space for grid search
search_space = [
    {
        'clf__n_estimators': [100, 300],
        'clf__learning_rate': [0.01, 0.1],
        'clf__max_depth': range(2, 5),
        'clf__colsample_bytree': [i/10.0 for i in range(4, 8)],    #change range (4,7) >> (4,8)
        'clf__gamma': [i/10.0 for i in range(3)],
        'fs__score_func': [chi2],
        'fs__k': [2],
    }
]
# Define cross validation
kfold = KFold(n_splits=5, random_state=60)    #Random state from 42 >> 60
# AUC and accuracy as score
scoring = {'AUC': 'roc_auc', 'Accuracy': make_scorer(accuracy_score), 'F1 score': 'f1_micro'}
# Define grid search
grid = GridSearchCV(
    pipe,
    param_grid=search_space,
    cv=kfold,
    scoring=scoring,
    refit='AUC',
    verbose=1,
    n_jobs=-1
)

# Fit grid search
xgb_model_clv_GS = grid.fit(X_train_clv, y_train_clv)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_split.py:296: FutureWarning:
FutureWarning
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
Fitting 5 folds for each of 144 candidates, totalling 720 fits
[Parallel(n_jobs=-1)]: Done 88 tasks    | elapsed: 15.5s
[Parallel(n_jobs=-1)]: Done 388 tasks  | elapsed: 1.1min
[Parallel(n_jobs=-1)]: Done 720 out of 720 | elapsed: 2.1min finished
```

```
[95] predicted_y = []
expected_y = []
predictions = xgb_model_clv_GS.predict(X_test_clv)
print('Best AUC Score: {}'.format(xgb_model_clv_GS.best_score_))
print('Accuracy: {}'.format(accuracy_score(y_test_clv, predictions)))
print(confusion_matrix(y_test_clv, predictions))

predicted_y.extend(predictions)
expected_y.extend(y_test_clv)
report_test = classification_report(expected_y, predicted_y)
print('test set')
print(report_test)
```

Best AUC Score: 0.7089054675375633

Accuracy: 0.6079380445304937

[[1089 759]

[51 167]]

test set

	precision	recall	f1-score	support
0	0.96	0.59	0.73	1848
1	0.18	0.77	0.29	218
accuracy			0.61	2066
macro avg	0.57	0.68	0.51	2066
weighted avg	0.87	0.61	0.68	2066